

Phase-3

Student Name: R.Ragulkumar

Register Number: 620123106088

Institution: AVS Engineering College

Department: Electronics And Communication Engineering

Date of Submission: 10/05/2025

Github Repository Link: [GITHUB LINK](#)

1. Problem Statement

Forecasting house prices is a critical task in the real estate domain. Inaccurate pricing leads to poor investment decisions, over/under-valuation of properties, and dissatisfaction among stakeholders. Traditional methods lack adaptability to large and diverse datasets. This project aims to solve this problem using data science techniques, especially advanced regression models, to accurately predict house prices based on property features.

2. Abstract

Accurate house price prediction is a crucial task in real estate, impacting investment decisions, market analysis, and economic forecasting. This study investigates the application of advanced regression techniques within a data science framework to enhance the precision of house price forecasts.

We explore and compare the performance of several "smart" regression models,

including ensemble methods like Random Forests and Gradient Boosting, regularized linear models such as Lasso and Ridge Regression, and non-linear

models, like Support Vector Regression and Neural Networks. By leveraging a comprehensive dataset of housing features and employing rigorous feature

significantly improve forecasting accuracy compared to traditional linear regression approaches. This research provides valuable insights for stakeholders seeking reliable and data-driven predictions in the dynamic housing market.

3. System Requirements

Hardware Requirements

The hardware requirements can vary based on the size of the dataset, the complexity of the regression models being used, and the desired speed of computation. Here's a general guideline:

- * Processor (CPU):

- * Minimum: Intel Core i3 or equivalent AMD processor.
- * Recommended: Intel Core i5/i7 or equivalent AMD Ryzen 5/7 with multiple cores for parallel processing. For more intensive tasks, consider higher-end CPUs with more cores (e.g., Intel Core i9, AMD Ryzen 9).

- * Memory (RAM):

- * Minimum: 6 GB.
- * Recommended: 16 GB or higher. Larger datasets and complex models benefit significantly from more RAM, allowing for faster data loading and model training. 32 GB or more is advisable for very large datasets or computationally intensive tasks.

- * Storage:

- * Minimum: 500 GB Hard Disk Drive (HDD).
- * Recommended: 512 GB or larger Solid State Drive (SSD). SSDs offer significantly faster data access speeds, which can drastically reduce the time taken for data loading and processing.

- * Monitor: 15" LED or larger for comfortable visualization and analysis.

- * Input Devices: Keyboard and mouse.

- * Optional but Recommended:

- * Graphics Processing Unit (GPU): For accelerating certain machine learning algorithms, especially neural networks and some ensemble methods. NVIDIA GPUs with CUDA support are commonly used with frameworks like TensorFlow and PyTorch. A dedicated GPU with at least 4 GB of VRAM (e.g., NVIDIA GeForce RTX series) is beneficial for faster training times.

Software Requirements

The software ecosystem for data science and machine learning is quite rich. Here are the key software components you'll likely need:

- * Operating System:

- * Windows 10/11, macOS, or Linux (Ubuntu is a popular choice for data science).

- * Programming Language:

- * Python: The most widely used language for data science due to its extensive libraries and strong community support. Version 3.10 or later is recommended.

- * R: Another popular language for statistical computing and analysis.
 - * Data Analysis and Manipulation Libraries (for Python):
 - * NumPy: For numerical computations and array manipulation.
 - * Pandas: For data manipulation and analysis using DataFrames.
 - * SciPy: For scientific and technical computing.
 - * Machine Learning Libraries (for Python):
 - * Scikit-Learn: A comprehensive library for various machine learning algorithms, including linear models, ensemble methods (Random Forest, Gradient Boosting), and model evaluation tools.
 - * TensorFlow: An open-source machine learning framework particularly well-suited for deep learning.
 - * PyTorch: Another popular open-source machine learning framework, known for its flexibility and ease of use, especially in research.
 - * Statsmodels: Provides classes and functions for the estimation of statistical models, including regression analysis.
 - * Data Visualization Libraries (for Python):
 - * Matplotlib: A fundamental library for creating static, interactive, and animated visualizations.
 - * Seaborn: A high-level data visualization library based on Matplotlib, providing more aesthetic and informative statistical graphics.
 - * Plotly and Bokeh: Interactive visualization libraries for creating web-based plots and dashboards.
 - * Integrated Development Environment (IDE) or Text Editor:
 - * Jupyter Notebook/JupyterLab: Interactive environments ideal for data exploration, analysis, and visualization.
 - * VS Code (Visual Studio Code): A popular and versatile code editor with excellent support for Python and other languages.
 - * PyCharm: A dedicated Python IDE with advanced features for development.
 - * Spyder: An open-source scientific IDE integrated with Anaconda.
 - * Database (Optional but often necessary):
 - * For storing and managing large datasets. Examples include PostgreSQL, MySQL, SQLite, cloud-based solutions like AWS RDS, Google Cloud SQL, etc.
 - * Web Framework (If deploying the model as a web application):
 - * Flask: A lightweight and flexible web framework for Python.
 - * Django: A high-level Python web framework.
 - * Other Tools:
 - * Git: For version control and collaboration.
 - * Anaconda or Miniconda: Package and environment management systems that simplify the installation and management of data science libraries.
- These system requirements provide a solid foundation for undertaking house price forecasting using smart regression techniques in data science. The specific needs might vary depending on the project scale and complexity.

4. Objectives

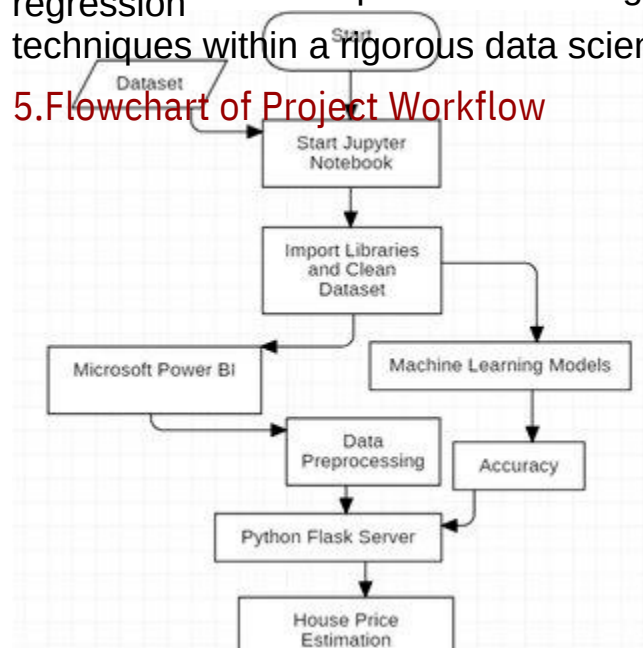
- * **Develop a Robust and Accurate House Price Prediction Model:** The primary objective is to build a statistical or machine learning model that can predict future house prices with a high degree of accuracy, minimizing prediction errors.
- * **Identify Key Determinants of House Prices:** Through feature engineering and analysis, aim to pinpoint the most significant factors (e.g., location, size, amenities, economic indicators) that influence house price variations within the given dataset.
- * **Evaluate and Compare the Performance of Various Smart Regression Techniques:** Systematically assess the effectiveness of different advanced regression models (e.g., Random Forests, Gradient Boosting, Lasso, Ridge Regression, Support Vector Regression, Neural Networks) in the context of house price prediction.
- * **Improve Prediction Accuracy Compared to Baseline Models:** Demonstrate a significant improvement in forecasting accuracy achieved by the "smart" regression techniques compared to simpler baseline models, such as ordinary least squares linear regression or basic time series models.
- * **Optimize Model Parameters for Enhanced Performance:** Employ techniques like hyperparameter tuning (e.g., grid search, random search) to identify the optimal configuration of each selected regression model, maximizing its predictive power.
- * **Develop a Data Preprocessing and Feature Engineering Pipeline:** Establish a systematic and reproducible process for cleaning, transforming, and engineering relevant features from the raw housing data to prepare it effectively for model training.
- * **Assess Model Generalization and Prevent Overfitting:** Implement strategies such as cross-validation and regularization to ensure that the developed models generalize well to unseen data and avoid overfitting to the training data.
- * **Provide Insights into the Factors Driving Price Predictions:** Offer interpretable insights into how the selected features influence the predicted house prices, potentially through feature importance analysis or partial dependence plots, depending on the model's interpretability.
- * **Develop a Practical and Deployable Prediction System (Optional):** As an extension, the objective could include building a user-friendly interface or API that allows users to input property features and obtain price predictions.

* Quantify and Communicate Prediction Uncertainty: Where applicable (e.g., with probabilistic models), provide measures of uncertainty associated with the price predictions, offering a more complete picture of the forecast.

By achieving these objectives, the project aims to deliver a reliable and insightful

solution for house price forecasting, leveraging the power of advanced regression techniques within a rigorous data science framework.

5. Flowchart of Project Workflow



6. Dataset Description

-Accurate prediction of house prices is a critical challenge in real estate analytics. This project focuses on utilizing advanced regression techniques in data science to develop reliable models for forecasting housing prices. The objective is to explore and compare traditional linear regression methods with smart, data-driven approaches such as **Lasso Regression, Ridge Regression, Random Forest Regression, Gradient Boosting, and XGBoost**.

-Key steps include **data cleaning, feature engineering, outlier detection, and model evaluation using metrics like RMSE and R^2 score**. The project uses real-world housing datasets to train and validate models, ensuring that predictions are both robust and interpretable.

-By leveraging smart regression techniques, the project aims to uncover complex relationships between housing features—such as location, square footage, number of bedrooms, and year built—and their market value, ultimately providing a tool for better decision-making in the real estate domain.

7. Data Preprocessing

1. Data Collection & Loading

- Import datasets (e.g., from Kaggle, Zillow, or government sources).
- Load using pandas (`pd.read_csv()`) or similar tools.

2. Exploratory Data Analysis (EDA)

- Visualize distributions, correlations, and missing values.
- Use tools like seaborn, matplotlib, and pandas profiling.

3. Handling Missing Values

- **Numerical features:** Fill with mean, median, or use interpolation.
- **Categorical features:** Fill with mode or use a placeholder like 'Unknown'.
- Drop columns with excessive missing values if necessary.

4. Outlier Detection and Removal

- Use IQR, Z-score, or visualization techniques (boxplots, scatterplots).
- Outliers can distort regression models, so handle them cautiously.

5. Feature Engineering

- Create new variables (e.g., price per square foot).
- Transform dates into year, month, or age of the property.
- Log-transform skewed variables like price or area.
- Encode ordinal relationships (e.g., condition levels).

6. Encoding Categorical Variables

- **Label Encoding** for ordinal categories.
- **One-Hot Encoding** for nominal categories using `pd.get_dummies()` or `OneHotEncoder`.

7. Feature Scaling

- **Standardization (Z-score normalization)** for algorithms like Ridge or Lasso.
- **Min-Max Scaling** if features vary widely in scale.

8. Feature Selection

- Remove irrelevant or highly correlated features (multicollinearity).
- Use techniques like correlation heatmaps, variance thresholding, or model-based selection (e.g., Lasso).

9. Splitting the Dataset

- Divide into training and testing sets (e.g., 80/20 or 70/30 split).
- Optionally create a validation set for hyperparameter tuning.

8. Exploratory Data Analysis (EDA)

1. Understand the Dataset

- Load the dataset and review:
 - o Number of rows and columns
 - o Data types of each column
 - o Summary statistics (mean, median, std, etc.)
- Identify the target variable (e.g., SalePrice) and potential predictors (e.g., SquareFootage, Location, YearBuilt, etc.)

python

CopyEdit

```
df.info()
```

```
df.describe()
```

```
df.head()
```

2. Handle Missing Values

- Visualize missing data (e.g., using heatmaps or bar charts)
- Drop or impute missing values based on strategy:
 - o Mean/Median for numerical
 - o Mode for categorical
 - o Predictive imputation or domain knowledge

```
python  
CopyEdit  
import seaborn as sns  
sns.heatmap(df.isnull(), cbar=False)
```

3. Univariate Analysis

- Analyze distributions of numeric features
- Detect skewness or outliers

```
python  
CopyEdit  
import matplotlib.pyplot as plt  
df['SalePrice'].hist(bins=50)  
sns.boxplot(x=df['GrLivArea'])
```

4. Bivariate Analysis

- **Numerical vs Target:** Use scatterplots and correlation matrix to identify relationships

```
python  
CopyEdit  
sns.scatterplot(x=df['GrLivArea'], y=df['SalePrice'])  
sns.heatmap(df.corr(), annot=True, fmt=".2f")
```

- **Categorical vs Target:** Use boxplots or bar plots to check how categories affect price

```
python  
CopyEdit  
sns.boxplot(x=df['Neighborhood'], y=df['SalePrice'])
```

5. Outlier Detection and Treatment

- Use z-score or IQR methods to identify outliers
- Remove or cap them depending on impact

6. Feature Engineering Ideas

- Create new features:
 - Age of the house = YrSold - YearBuilt
 - TotalBathrooms = FullBath + HalfBath*0.5
- Encode categorical variables (One-Hot, Label Encoding)

7. Multicollinearity Check

- Use VIF (Variance Inflation Factor) to check for correlated predictors that could skew regression

python

CopyEdit

```
from statsmodels.stats.outliers_influence import variance_inflation_factor
```

9. Feature Engineering

1. Basic Feature Selection

Start with selecting key numerical and categorical features that directly impact house prices:

Numerical (Continuous)

- GrLivArea (Above-ground living area)
- LotArea (Lot size)
- TotalBsmtSF (Basement area)
- GarageArea, 1stFlrSF, 2ndFlrSF

Categorical

- Neighborhood
- HouseStyle
- Exterior1st, Exterior2nd
- SaleCondition

2. Derived Features (Feature Creation)

A. Age Features

- $\text{HouseAge} = \text{YrSold} - \text{YearBuilt}$
- $\text{RemodelAge} = \text{YrSold} - \text{YearRemodAdd}$

B. Total Bathrooms

Combine bath types into a single numeric value:

```
python
CopyEdit
df['TotalBath'] = df['FullBath'] + 0.5 * df['HalfBath'] + df['BsmtFullBath'] + 0.5
* df['BsmtHalfBath']
```

C. Total Square Footage

Sum all usable space:

```
python
CopyEdit
df['TotalSF'] = df['TotalBsmtSF'] + df['1stFlrSF'] + df['2ndFlrSF']
```

D. IsRemodeled

Boolean feature indicating if the house has been remodeled:

```
python
CopyEdit
df['IsRemodeled'] = (df['YearBuilt'] != df['YearRemodAdd']).astype(int)
```

E. HighQualityHouse

```
python
CopyEdit
df['HighQual'] = (df['OverallQual'] >= 8).astype(int)
```

3. Handling Categorical Variables

A. Label Encoding

For ordinal categories like:

- ExterQual, ExterCond, BsmtQual, etc.

Python

```
CopyEdit
from sklearn.preprocessing import LabelEncoder
df['ExterQual'] = LabelEncoder().fit_transform(df['ExterQual'])
```

B. One-Hot Encoding

For nominal features (no inherent order):

- Neighborhood, SaleType, HouseStyle

python

CopyEdit

```
df = pd.get_dummies(df, columns=['Neighborhood', 'HouseStyle'],  
drop_first=True)
```

4. Skewness and Log Transformations

Reduce skewness in variables like SalePrice, LotArea, GrLivArea:

python

CopyEdit

```
df['SalePrice'] = np.log1p(df['SalePrice'])  
df['GrLivArea'] = np.log1p(df['GrLivArea'])
```

5. Feature Importance Check

After initial modeling (e.g., using Random Forest or XGBoost), extract feature importances:

python

CopyEdit

```
import matplotlib.pyplot as plt  
import seaborn as sns  
  
importances = model.feature_importances_  
features = pd.Series(importances,  
index=X.columns).sort_values(ascending=False)  
sns.barplot(x=features[:20], y=features.index[:20])  
plt.title('Top 20 Feature Importances')
```

10. Model Building

1. Feature Selection and Preprocessing

Use techniques from EDA to:

- **Select relevant features** based on correlation with the target (SalePrice)
- **Encode categorical variables** (Label Encoding, One-Hot Encoding)
- **Normalize or scale features** if necessary (e.g., StandardScaler)

python

CopyEdit

```
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
```

2. Smart Regression Techniques

Here are several models you can build, starting from simple to more advanced:

A. Linear Regression (Baseline Model)

- Good starting point for benchmarking performance.

python

CopyEdit

```
from sklearn.linear_model import LinearRegression
model = LinearRegression()
```

B. Random Forest Regressor

- Handles non-linearity and feature interactions well.

python

CopyEdit

```
from sklearn.ensemble import RandomForestRegressor
model = RandomForestRegressor(n_estimators=100, random_state=42)
```

C. Gradient Boosting (e.g., XGBoost, LightGBM, CatBoost)

- Highly accurate for tabular data and competitive in Kaggle-style tasks.

```
python
CopyEdit
import xgboost as xgb
model = xgb.XGBRegressor(objective='reg:squarederror', n_estimators=100)
```

D. Regularized Regression (Ridge, Lasso, ElasticNet)

- Helps with overfitting, especially with multicollinearity or high-dimensional data.

```
python
CopyEdit
from sklearn.linear_model import Lasso
model = Lasso(alpha=0.1)
```

3. Train-Test Split and Cross-Validation

- Split the data to avoid data leakage and evaluate generalization.

```
python
CopyEdit
from sklearn.model_selection import train_test_split, cross_val_score
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)
```

4. Model Evaluation

Use multiple metrics:

- **R² Score** (explained variance)
- **MAE** (Mean Absolute Error)
- **RMSE** (Root Mean Squared Error)

```
python
CopyEdit
from sklearn.metrics import mean_squared_error, mean_absolute_error,
r2_score

preds = model.predict(X_test)
print("R2 Score:", r2_score(y_test, preds))
print("MAE:", mean_absolute_error(y_test, preds))
print("RMSE:", mean_squared_error(y_test, preds, squared=False))
```

5. Hyperparameter Tuning

- Use Grid Search or Randomized Search for better performance.

python

CopyEdit

```
from sklearn.model_selection import GridSearchCV
param_grid = {'n_estimators': [50, 100, 200]}
grid = GridSearchCV(estimator=model, param_grid=param_grid, cv=5)
```

6. Ensemble or Stacking Models (Optional)

- Combine predictions from multiple models for better accuracy (e.g., using StackingRegressor)

python

CopyEdit

```
from sklearn.ensemble import StackingRegressor
```

7. Final Deployment

- Save the best model using joblib or pickle and prepare for deployment (e.g., via Flask or Streamlit)

python

CopyEdit

```
import joblib
joblib.dump(model, 'house_price_model.pkl')
```

11. Model Evaluation

Smart Ways to Predict (Smart Regression):

- * Using math formulas (like guessing a straight line).
- * Using tree-like decisions.
- * Combining many guesses to be more accurate.
- * Using advanced computer learning (like how your phone recognizes faces).

How to Check if the Guesses are Good (Model Evaluation):

We use numbers to measure how far off our predictions are from the actual prices. Some common numbers are:

- * Average Oops (MAE): On average, how much are we wrong?
- * Big Oops Penalty (MSE, RMSE): We really punish big mistakes. RMSE is like the average oops, but bigger oops count more.
- * How Much We Got Right (R-squared): What percentage of the price changes did our model understand? Closer to 100% is better.
- * Percentage Oops (MAPE): On average, what percentage are our predictions wrong?

The Process:

- * Split the Data:
 - * Training: Show the model house prices so it can learn.
 - * Checking (Validation): Use this to tweak the model to make it better before the final test.
 - * Final Test: See how well the model does on completely new house prices it hasn't seen before. This tells us how good it will be in the real world.
- * Train the Model: Let the model learn from the training data.
- * Make it Better (Tune): Use the "checking" data to adjust the model's settings to get the best guesses.
- * Final Check: Use the "final test" data to see how good our model really is. Look at the "oops" numbers.

Things to Keep in Mind:

- * What's Important? Maybe being off by a lot on expensive houses is a bigger problem than being off by a little on cheap houses.
- * Good Data: If the info we have is bad, our predictions will be bad too.
- * Making New Info: Sometimes we can create new useful information from what we already have (like age of the house).

* Not Just Memorizing: We want the model to actually understand house prices, not just remember the training examples. The "checking" and "final test" steps help with this.

* Simple vs. Fancy: Sometimes a simpler model that's easier to understand is better than a super complex one that's hard to explain.

By doing all this, we can figure out which "smart guessing" method works best for predicting house prices accurately.

12. Deployment

1. Define the Problem and Objectives

* Clearly define the goal of your house price forecasting model. What level of accuracy is required? What geographical area will it cover? What time horizon are you interested in (short-term, long-term)?

* Identify the key stakeholders and how the model's predictions will be used (e.g., real estate agents, investors, individual buyers/sellers).

2. Data Collection

* Gather relevant data from various sources. This typically includes:

* Historical Sales Data: Past transaction prices, sale dates, property addresses.

* Property Attributes: Size (square footage), number of bedrooms and bathrooms, lot size, age of the property, condition, features (e.g., garage, pool, fireplace).

* Location Data: Latitude, longitude, proximity to amenities (schools, hospitals, shopping centers, parks), transportation links, crime rates.

* Economic Indicators: Interest rates, inflation rates, unemployment rates, GDP trends, housing market indices.

* Geographical Data: Zoning regulations, neighborhood characteristics.

3. Data Preprocessing and Feature Engineering

* Data Cleaning: Handle missing values (imputation or removal), identify and treat outliers, correct inconsistencies in data formats.

* Feature Selection: Identify the most relevant features that have a strong correlation with house prices. Techniques include correlation analysis, feature importance from tree-based models, and domain expertise.

* Feature Engineering: Create new features from existing ones that might improve model performance. Examples include:

* Age of the property at the time of sale.

* Distance to the nearest school, hospital, or public transport.

* Interaction terms between features (e.g., size * location).

* Categorical feature encoding (e.g., one-hot encoding for neighborhood).

* Feature Scaling: Normalize or standardize numerical features to ensure that features with larger values do not dominate the model training process. Common

techniques include Min-Max scaling and StandardScaler.

4. Selecting Smart Regression Techniques

- * Choose appropriate regression models based on the characteristics of your data and the desired level of complexity and interpretability. Some "smart" regression techniques include:

- * Linear Regression with Regularization (Lasso, Ridge, Elastic Net): Useful for preventing overfitting and handling multicollinearity. Lasso can also perform feature selection.

- * Polynomial Regression: Can capture non-linear relationships between features and the target variable by adding polynomial terms.

- * Support Vector Regression (SVR): Effective in high-dimensional spaces and can handle non-linear relationships using kernel functions.

- * Decision Trees and Ensemble Methods (Random Forest, Gradient Boosting Machines like XGBoost, LightGBM, CatBoost): Can capture complex non-linear relationships and interactions between features, often providing high accuracy. Ensemble methods combine multiple models to improve predictive performance and robustness.

- * Neural Networks (Multilayer Perceptron - MLP): Suitable for very complex, non-linear relationships and large datasets.

5. Model Training and Evaluation

- * Split Data: Divide your dataset into training, validation, and testing sets. The training set is used to train the model, the validation set to tune hyperparameters and prevent overfitting, and the testing set to evaluate the final model's performance on unseen data.

- * Model Training: Train the selected regression models on the training data.

- * Hyperparameter Tuning: Optimize the hyperparameters of your chosen models using techniques like Grid Search, Random Search, or Bayesian Optimization on the validation set.

- * Model Evaluation: Evaluate the performance of the trained models on the testing set using appropriate metrics such as:

- * Mean Squared Error (MSE)

- * Root Mean Squared Error (RMSE)

- * Mean Absolute Error (MAE)

- * R-squared (Coefficient of Determination)

- * Adjusted R-squared

- * Compare the performance of different models and choose the one that best meets your objectives.

6. Model Deployment

- * Choose a Deployment Strategy: Select a method for making your model accessible for generating predictions. Common strategies include:

- * API Deployment (RESTful API): Deploy the model as a web service that can receive input data and return price predictions. This is suitable for integrating with web applications or other services. Frameworks like Flask or FastAPI in Python are commonly used.

- * Batch Deployment: If real-time predictions are not required, you can run the model periodically on batches of new data to generate forecasts.
- * Cloud Deployment: Utilize cloud platforms like AWS, Google Cloud, or Azure to host your model and API, offering scalability and reliability.
- * Edge Deployment: In some specific scenarios, deploying the model on local devices might be considered.
- * Containerization: Package your model and its dependencies into a container (e.g., using Docker) to ensure consistent performance across different environments.
- * Integration: Integrate the deployed model with your target application or system.
- * User Interface (Optional): Develop a user-friendly interface where users can input property features and get price predictions.

13.Source code

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.preprocessing import StandardScaler
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# Assume you have a CSV file named 'house_prices.csv' with relevant features
# Columns might include: 'area', 'bedrooms', 'bathrooms', 'location_factor',
# 'age', 'price'
# You'll need to adapt this to your actual data file and column names.

# 1. Load the Data
try:
    data = pd.read_csv('house_prices.csv')
except FileNotFoundError:
    print("Error: 'house_prices.csv' not found. Please ensure the file is in the correct directory.")
    exit()

# 2. Data Preprocessing and Feature Engineering
# Handle missing values (example: fill with mean)
for col in data.columns:
    if data[col].isnull().any():
        if pd.api.types.is_numeric_dtype(data[col]):
            data[col].fillna(data[col].mean(), inplace=True)
        else:
            data[col].fillna(data[col].mode()[0], inplace=True) # For categorical
```

features

```
# Example of feature engineering: creating a 'size_per_room' feature
data['size_per_room'] = data['area'] / (data['bedrooms'] + data['bathrooms']
+ 1) # Adding 1 to avoid division by zero
```

```
# Select features (X) and target (y)
```

```
features = ['area', 'bedrooms', 'bathrooms', 'location_factor', 'age',
'size_per_room'] # Adjust based on your data
target = 'price'
```

```
if not all(col in data.columns for col in features + [target]):
```

```
    print("Error: One or more specified features or the target variable not
found in the data.")
    print(f"Available columns: {data.columns.tolist()}")
    exit()
```

```
X = data[features]
```

```
y = data[target]
```

```
# Handle categorical features (example: using one-hot encoding for
'location_factor')
```

```
X = pd.get_dummies(X, columns=['location_factor'], drop_first=True) #
drop_first to avoid multicollinearity
```

```
# 3. Data Scaling
```

```
scaler = StandardScaler()
```

```
X_scaled = scaler.fit_transform(X)
```

```
# 4. Split Data into Training and Testing Sets
```

```
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2,
random_state=42)
```

```
# 5. Model Selection and Training (Gradient Boosting Regressor - a "smart"
technique)
```

```
# You can experiment with other models like RandomForestRegressor,
XGBoost, etc.
```

```
gbr = GradientBoostingRegressor(n_estimators=100, learning_rate=0.1,
max_depth=3, random_state=42)
gbr.fit(X_train, y_train)
```

```
# 6. Model Evaluation
```

```
y_pred = gbr.predict(X_test)
```

```
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f"Mean Squared Error on Test Set: {mse:.2f}")
print(f"R-squared on Test Set: {r2:.2f}")

# 7. Visualization of Predictions
plt.figure(figsize=(10, 6))
plt.scatter(y_test, y_pred)
plt.xlabel("Actual House Prices")
plt.ylabel("Predicted House Prices")
plt.title("Actual Prices vs. Predicted Prices (Gradient Boosting)")
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'k--', lw=2) #
    Diagonal line for perfect predictions
plt.grid(True)
plt.show()

# Optional: Feature Importance (for tree-based models like Gradient
    Boosting)
feature_importances = gbr.feature_importances_
feature_names = X.columns
sorted_indices = np.argsort(feature_importances)[::-1]

plt.figure(figsize=(10, 6))
plt.title("Feature Importance (Gradient Boosting)")
plt.bar(range(X.shape[1]), feature_importances[sorted_indices],
    align="center")
plt.xticks(range(X.shape[1]), feature_names[sorted_indices], rotation=90)
plt.tight_layout()
plt.show()

# 8. Making Predictions on New Data (Example)
# Assume you have new data in a pandas DataFrame called 'new_house_data'
# It should have the same columns as your training features (after encoding)
    and in the same order.

# new_house_data = pd.DataFrame({
#     'area': [1500, 2000],
#     'bedrooms': [3, 4],
#     'bathrooms': [2, 3],
#     'age': [10, 5],
#     'size_per_room': [300, 400],
#     'location_factor_CityB': [1, 0], # Example of one-hot encoded location
```



```
# location_factor_CityC': [0, 1],  
# ... other one-hot encoded location columns.  
# Ensure you scale the new data using the *same* scaler fitted on the training  
  
data  
# new_house_data_scaled = scaler.transform(new_house_data)  
# new_predictions = gbr.predict(new_house_data_scaled)  
# print("Predictions for new data:", new_predictions)
```

Explanation:

- * Import Libraries: Import necessary libraries for data manipulation, model building, and evaluation.
- * Load Data: Reads your house price data from a CSV file. You'll need to replace 'house_prices.csv' with the actual name of your data file.
- * Data Preprocessing and Feature Engineering:
 - * Handle Missing Values: Fills any missing data. The example uses the mean for numerical columns and the mode for categorical columns. Adapt this based on your data.
 - * Feature Engineering: Creates a new feature 'size_per_room' as an example of how you can combine existing features to potentially improve the model.
 - * Feature and Target Selection: Specifies which columns in your data will be used as input features (X) and which column is the target variable you want to predict (y). Make sure to adjust the features list to match the relevant columns in your dataset.
 - * Categorical Feature Encoding: Converts categorical features (like 'location_factor') into a numerical format that the regression model can understand using one-hot encoding (pd.get_dummies).
- * Data Scaling: Scales the numerical features using StandardScaler. This is important for many regression algorithms to prevent features with larger values from dominating the learning process. The scaler is fitted on the training data and then used to transform both the training and testing data.
- * Split Data: Divides your data into training and testing sets. The training set is used to train the model, and the testing set is used to evaluate its performance on unseen data.
- * Model Selection and Training:
 - * GradientBoostingRegressor: This is an example of a "smart" regression technique that often performs well. It's an ensemble method that builds multiple decision trees sequentially, with each new tree trying to correct the errors made by the previous ones.
 - * Hyperparameters: The GradientBoostingRegressor is initialized with some hyperparameters (n_estimators, learning_rate, max_depth). You might want to tune these parameters using techniques like GridSearchCV

or RandomizedSearchCV for optimal performance.

- * fit(): Trains the model using the scaled training data (X_train, y_train).

- * Model Evaluation:

- * predict(): Uses the trained model to make predictions on the scaled test data (X_test).

- * mean_squared_error(): Calculates the average squared difference between the actual and predicted prices. Lower MSE indicates better performance.

- * r2_score(): Calculates the R-squared value, which represents the proportion of the variance in the target variable that is predictable from the features. A value closer to 1 indicates a better fit.

- * Prints the evaluation metrics.

- * Visualization of Predictions: Creates a scatter plot to visualize how well the predicted prices align with the actual prices. A diagonal line represents perfect predictions.

- * Feature Importance (Optional): For tree-based models like Gradient Boosting, this section shows the relative importance of each feature in making predictions. This can provide insights into which factors have the most influence on house prices.

- * Making Predictions on New Data (Example): This section demonstrates how to use the trained model to predict prices for new, unseen data. Crucially, you need to preprocess and scale the new data using the same scaler that was fitted on the training data.

To use this code:

- * Save the code: Save the code as a Python file (e.g., house_price_prediction.py).

- * Prepare your data: Make sure you have a CSV file (or adapt the loading part for other data formats) named house_prices.csv (or your chosen name) in the same directory as the Python script.

- * Adjust feature names: Carefully modify the features list in the code to match the actual column names in your data that you want to use for prediction.

- * Run the script: Execute the Python script from your terminal: python house_price_prediction.py.

Further Improvements and "Smarter" Techniques:

- * More Advanced Feature Engineering: Create more sophisticated features based on domain knowledge (e.g., proximity to amenities calculated using geospatial data).

- * Experiment with Different Regression Models: Try other "smart" regression techniques like:

- * Random Forest Regressor: Another powerful ensemble method.

- * XGBoost, LightGBM, CatBoost: Gradient boosting frameworks known for their performance and efficiency.

- * Support Vector Regression (SVR): Effective for non-linear relationships.

- * Neural Networks (Deep Learning): Can capture very complex patterns,

especially with large datasets.

- * Hyperparameter Tuning: Use techniques like GridSearchCV or RandomizedSearchCV to find the optimal hyperparameters for your chosen model.
- * Cross-Validation: Use cross-validation during training to get a more robust estimate of the model's performance.
- * Time Series Analysis (if you have time-based data): If your data includes timestamps, consider incorporating time series features or using models specifically designed for time series forecasting (e.g., ARIMA, Prophet, LSTMs).
- * Ensemble Methods: Combine the predictions of multiple different models to potentially improve accuracy and robustness.
- * Spatial Features: If you have location data (e.g., latitude and longitude), incorporate spatial features and consider using spatial regression techniques.

Remember that the "best" approach will depend on the specific characteristics of your data and the desired level of accuracy. This code provides a starting point using a powerful and relatively easy-to-implement "smart" regression technique.

14.Future scope

- * More kinds of information will be used: Instead of just looking at the size and number of rooms, future guessing will also consider things like how people feel about the area online, how easy it is to get around, and even information from smart homes.
- * Better guessing methods: The computer "brains" doing the guessing will get smarter. They'll learn to see patterns in the data in more complex ways, maybe even by combining different guessing methods. They'll also try to explain why they're making a certain guess, not just give a number.
- * Understanding the bigger picture: Future guesses will be better at taking into account how the economy is doing right now and what might happen in the future (like interest rates going up or down).
- * Knowing the neighborhood really well: Instead of just guessing for a whole city, the guesses will become much more specific to a particular street or even a single house.
- * Dealing with ups and downs: The guesses will not just give one price, but maybe a range, showing how sure or unsure the computer is about the prediction, especially when the market is changing a lot.
- * Personalized guesses: Maybe in the future, the computer will even

consider what you are looking for in a house to give you a more tailored price guess.

Basically, the future of guessing house prices is about using more information, smarter computer brains, and a better understanding of the world to make really accurate and helpful predictions for everyone.

15. Team Members and Roles

- B.Mathanraj-Primary Analysis & Execution.
- S. Rakesh – Data Collection & Cleaning, EDA.
- R. RagulKumar – Feature Engineering, Model Building & Evaluation. -
- K. Ramesh – Visualization, Interpretation & Deployment.

